

# quiz

April 13, 2020

```
In [11]: from tensorflow.examples.tutorials.mnist import input_data
import tensorflow as tf
import numpy as np
from helper import batches

learning_rate = 0.001
n_input = 784 # MNIST data input (img shape: 28*28)
n_classes = 10 # MNIST total classes (0-9 digits)

# Import MNIST data
mnist = input_data.read_data_sets('/datasets/ud730/mnist', one_hot=True)

# The features are already scaled and the data is shuffled
train_features = mnist.train.images
test_features = mnist.test.images

train_labels = mnist.train.labels.astype(np.float32)
test_labels = mnist.test.labels.astype(np.float32)

# Features and Labels
features = tf.placeholder(tf.float32, [None, n_input])
labels = tf.placeholder(tf.float32, [None, n_classes])

# Weights & bias
weights = tf.Variable(tf.random_normal([n_input, n_classes]))
bias = tf.Variable(tf.random_normal([n_classes]))

# Logits -  $xW + b$ 
logits = tf.add(tf.matmul(features, weights), bias)

# Define loss and optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=labels))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cost)

# Calculate accuracy
correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(labels, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

```
Extracting /datasets/ud730/mnist/train-images-idx3-ubyte.gz
Extracting /datasets/ud730/mnist/train-labels-idx1-ubyte.gz
Extracting /datasets/ud730/mnist/t10k-images-idx3-ubyte.gz
Extracting /datasets/ud730/mnist/t10k-labels-idx1-ubyte.gz
```

```
In [37]: # TODO: Set batch size
        batch_size = 128
        assert batch_size is not None, 'You must set the batch size'

        init = tf.global_variables_initializer()

        with tf.Session() as sess:
            sess.run(init)

            # TODO: Train optimizer on all batches
            # for batch_features, batch_labels in -----
            for batch_features, batch_labels in batches(batch_size, train_features, train_labels):
                sess.run(optimizer, feed_dict={features: batch_features, labels: batch_labels})

            # Calculate accuracy for test dataset
            test_accuracy = sess.run(
                accuracy,
                feed_dict={features: test_features, labels: test_labels})

            print('Test Accuracy: {}'.format(test_accuracy))
```

```
Test Accuracy: 0.11150000244379044
```

```
In [ ]:
```